

PRACHI RAVI SAIBEWAR

Orlando, FL | Open to Remote & Relocation | +1 386-456-8740 | prachi.saibewar.cs@gmail.com | [LinkedIn](#) | [GitHub](#)

Summary

AI Software Engineer with **4+ years of software engineering experience**, including production Generative AI, LLM, RAG, and agentic AI systems. Developed AI-assisted workflows in regulated healthcare and enterprise environments, improving grounded-answer rates from **79% to 92%** and retrieval Recall@5 from **76% to 91%**, with experience across retrieval, evaluation, backend integration, observability, and production support.

Technical Skills

AI & Generative AI: Generative AI (GenAI), Machine Learning (ML), Large Language Models (LLMs), OpenAI, Vertex AI, Retrieval-Augmented Generation (RAG), Agentic AI, AI Agents, LangChain, LangGraph, NLP, Prompt & Context Engineering, Embeddings, Vector Search, Semantic Search, Chunking, Structured Outputs, Function/Tool Calling, LLM Evaluation, Retrieval Evaluation, Langfuse, Guardrails, Human-in-the-Loop

Languages & Backend: Python, C#, TypeScript, JavaScript, SQL, FastAPI, ASP.NET Core/.NET, Node.js, REST APIs, Microservices, Distributed Systems, System Design, Event-Driven Architecture, OAuth2

Cloud & Data: AWS, S3, SQS, PostgreSQL, pgvector, Redis, Docker, CI/CD, Git

Frontend, Testing & Observability: React, OpenTelemetry, Prometheus, PromQL, Chronosphere, Jest, Playwright, Postman

Experience

Software Engineer | Rokt | Boca Raton, FL | Sep 2025 – Aug 2026

- Built a Customer 360 identity-resolution investigation workflow using React, TypeScript, C#/.NET REST APIs, AWS S3, and indexed search, enabling engineers to trace profile merges and reducing investigation time by **30%**.
- Engineered consent-aware data routing across **300+ native Customer Data Platform (CDP) integrations** and UID2 identity workflows, propagating privacy and preference signals through downstream activation pipelines to support GDPR/CCPA-aligned processing.
- Implemented production observability with Prometheus, PromQL, and Chronosphere, introducing **80% and 95% schema-capacity thresholds** to identify ingestion and reliability risks before platform limits were reached.
- Automated recurring audience data exchanges using AWS S3 and SFTP, reducing manual operational effort by **80%** and improving reliability of scheduled downstream delivery.

AI Engineer | Orlando Health | Orlando, FL | Jan 2025 – Sep 2025

- Engineered retrieval-augmented generation (RAG) over 10K+ knowledge records with embeddings, chunking, metadata filtering, semantic vector search, and PostgreSQL/pgvector, improving Recall@5 from 76% to 91% across a 300-query evaluation set.
- Implemented structured-output LLM workflows with prompt/context engineering, validation, guardrails, and human-in-the-loop review, increasing **schema-valid responses to 98%+** and reducing manual corrections by **35%**.
- Established an automated LLM evaluation framework across **300+ queries**, improving grounded-answer rate from **79% to 92%** through retrieval and prompt iteration.
- Engineered event-driven AI/data pipelines using Python, PostgreSQL, AWS S3, and SQS, adding retries, logging, failure handling, and diagnostics while reducing **p95 processing latency by 35%**.
- Built REST-based healthcare integrations using **FHIR R4 and HL7 v2 (ADT/ORU/ORM)**, enabling secure exchange of patient and clinical data across EHR systems with **OAuth 2.0 and HIPAA-aligned access controls for PHI**.

Software Engineer | ServiceNow | Remote, India | May 2020 – Jul 2023

- Developed reusable enterprise workflow automation and REST integration services in JavaScript, translating business requirements into **20+ production workflows** and reducing manual processing by **40%**.
- Integrated **15+ internal and third-party systems** through REST APIs and asynchronous workflows, implementing transformation, validation, retries, and exception handling for reliable cross-system data exchange.
- Optimized backend and database paths through query tuning and caching, reducing API/workflow processing latency by **25%** across high-volume enterprise operations.

PROJECTS

INCIDENTIQ — AGENTIC AI PRODUCTION SUPPORT PLATFORM ([Github](#))

Python | FastAPI | React | TypeScript | LangGraph | RAG | PostgreSQL/pgvector | Langfuse | Docker

- Built a full-stack AI incident-investigation agent with LangGraph, RAG, pgvector, tool calling, retries, and stateful diagnostics.
- Preserved agent state and tool context across failures/retries, achieving 90%+ task completion across a 100-scenario agent evaluation suite.
- Instrumented agents with Langfuse tracing, LLM evaluation, and human-in-the-loop guardrails, achieving 90%+ grounded responses and 95%+ tool-execution success.

Education

- Master of Science in Computer Science | University of Central Florida | Orlando, FL | Aug 2023 – May 2025**
- Bachelor of Engineering in Computer Science | Savitribai Phule Pune University | Pune, India | Aug 2016 – May 2020**